

CS417 Lab #7

Getting Started

Begin the lab by creating a folder for your files. Then, go to `mycourses.unh.edu`, find Lab07, and download these files into that folder:

- `numbers.txt`
- `histogram.py`
- `roll_die.py`

Exercises

1. The file `numbers.txt` contains 1000 numbers. Modify `histogram.py` to compute a *histogram* of these numbers. A histogram consists of several counts, corresponding to several bins. Each count says how many values fall into one bin.

To run the program, type this on the command line:

```
python histogram.py numbers.txt 10
```

You should get this output:

```
22
41
112
152
192
189
148
82
52
10
```

I've already written code which reads the numbers from the file in `sys.argv[1]`, and gets `n_bins`, the number of histogram bins, from `sys.argv[2]`.

Find the min (`lo`) and max (`hi`) of the numbers. Then find the `span = hi - lo` of the numbers. Finally, increase `span` by a little bit:

`span *= 1.0001`, to make it bigger than `lo - hi`.

Now, split the `span` into equal-width intervals, or “bins”. There should be `n_bins` of them. Given a number `x`, which bin does it go into? Here's the arithmetic you'll need:

- x is between lo and hi , thus
- $x - lo$ is between 0 and $hi - lo$, thus
- $(x - lo) / span$ is between 0.0 and 1.0, thus
- $(x - lo) / span * n_bins$ is between 0.0 and n_bins , so write this:
- `bin = int((x - lo) / span * n_bins)`

Because `span` is a bit larger than $hi - lo$, the `bin` will never be `n_bins`. It will range from 0 upto `n_bins - 1`. This covers exactly `n_bins` possible different values, which is what we want.

2. Open `roll_die.py`, and implement the function `cummulative_probabilities(probs)`.

To do this, build the cumulative probabilities. This is a list of values. Each value is the sum of values in `probs`, thus:

```
cumm_probs[0] = probs(0)
cumm_probs[1] = probs(0) + probs(1)
cumm_probs[2] = probs(0) + probs(1) + probs(2)
...
```

However, you can't type that, because you don't know the length of `probs`; you need a more general approach. Do this:

- start with an empty list: `cumm_probs = []`.
 - the first `cumm_probs[0]` is `probs[0]` (just append that to `cumm_probs`)
 - all the other `cumm_probs[i]` are `cumm_probs[i - 1] + probs[i]` (append that). You'll need `for i in range(1, len(probs)):`
-

3. In `roll_die.py`, implement the generator `biased_generator()`. It is passed a list of cumulative probabilities, and generates a sequence of biased random values.

To do this, first call `x = random.random()` to generate a random number `x` from 0.0 to 1.0.

Visit all the cumulative probabilities, and stop when `x` exceeds it the first time:

```
for k, prob in enumerate(cumm_probs):
    if prob > x:
```

You want the `k` where this first occurs: break the `for` loop, and then `yield k`.

Notice that `biased_generator()` doesn't return; it yields. This works with the `for x` loop in the `main()` function, to produce a sequence of values.

4. Test your programs: the `main()` function runs the generator 10000 times, and prints the die values. Save these in a file, and then use `histogram.py` to count them. You can do it from the command line thus:

```
python roll_die.py > rolls.txt  
python histogram.py rolls.txt 6
```

The biased die has these probabilities:

```
0 0.1  
1 0.2  
2 0.1  
3 0.2  
4 0.15  
5 0.25
```

So, if your program works, it should generate 10,000 counts distributed *approximately* like this:

```
1000  
2000  
1000  
2000  
1500  
2500
```

Turning in your work

At the end of the lab session, turn in any work you have completed. You can continue to work today, and turn in all the work by midnight, with no lateness penalty. Go to `mycourses.unh.edu`, find lab 9, click “Submit”, and upload `histogram.py` and `roll_die.py`.