

**IT 775**  
**Database Technology**  
**SQL-DML**

**Select Statement**  
**SubQueries**

# Nested Subqueries

- inner query returns information to an outer query
- subquery is a parenthesized select statement embedded in another select statement
- subquery can occur in the WHERE clause of outer query

SELECT columns FROM tables

WHERE predicate

( SELECT columns FROM tables WHERE cond-expr )


- subquery can occur in the FROM clause of outer query

SELECT columns

FROM table JOIN ( SELECT cols FROM tbls WHERE pred )

WHERE predicate

- or both



# Rules for Subqueries

each subquery must be enclosed in parenthesis

ORDER BY clause can appear only at the outermost level

subqueries can nest to any level

subquery can reference relation tables and attribute names in containing (outer) queries

subquery cannot reference relation tables and attributes names in contained (inner) clauses

ambiguous unqualified attribute references (no relation name or range variable) refer to innermost inclusion level

see example next

# Subquery Example

SELECT \* FROM student S  
 WHERE NOT EXISTS  
 ( SELECT \* FROM enroll E  
 WHERE snbr = S.snbr  
 AND EXISTS  
 ( SELECT \* FROM prereq P  
 WHERE dept = E.dept AND cnbr = E.cnbr  
 AND NOT EXISTS  
 ( SELECT \* FROM transcript  
 WHERE dept = P.pdept AND cnbr = P.pcnbr  
 AND snbr = S.snbr AND grade >= 1.0  
 ) )  
 )  
 )  
 AND EXISTS  
 ( SELECT \* FROM enroll WHERE snbr = S.snbr  
 )

---> *ref within subquery*  
 ....> *ref containing (sub)query*  
 ———> *ref outer query*

# Subquery Result Forms & Usage

subquery can return:

one column of values (degenerate table)

- single row — **scalar** subquery
- multiple rows — **set-valued** subquery

multiple columns of values (full table)

- **existence** subquery only
  - returns true if at least one row qualifies
  - returns false if subquery result is empty
- set-valued subquery + row constructor
  - tuple is treated as an element in a set.

## subquery usage

FROM clause – arbitrary table result – aliased

WHERE/HAVING clause – acceptable result: f(comparator)

- = <> < > BETWEEN      scalar result only
- IN ALL ANY      set result only (single column)
- EXISTS, NOT EXISTS      full table result

# Scalar Subquery Result Usage

table of 1 column and 1 row — i.e., single value  
can be used directly in conditional expressions  
if value is ordinal, it can be an operand in a comparison expr  
Example: students whose gpa exceeds that of Joe College

```
SELECT sname FROM student
WHERE gpa > ( SELECT gpa FROM student
              WHERE sname = 'joe college'
            )
```

Example: students in the same major as Joe College

```
SELECT sname FROM student
WHERE major = ( SELECT major FROM student
                WHERE sname = 'joe college'
              )
```

# Statistical Subquery Scalar-Valued

- students with above average gpa values

- stat subquery version

```
SELECT sname, gpa FROM student  
WHERE gpa > ( SELECT avg( gpa ) FROM student );
```

- students with above average in-major gpa values

```
SELECT sname, major, gpa FROM student S  
WHERE gpa >  
  ( SELECT avg( gpa ) FROM student WHERE major = S.major );
```

note: horrible efficiency — this correlated subquery reevaluates the average gpa of a major for each student in that major

(unless optimizer rescues us)

# Set-Valued Subquery Result Usage

must be in the form of a table of 1 column  
usually more than 1 row — i.e., multi-valued  
used in set comparison expressions

expr  $\theta$  SOME (set-valued subquery)

expr  $\theta$  ANY (set-valued subquery)

expr  $\theta$  ALL (set-valued subquery)

expr [NOT] IN (set-valued subquery)

e.g., CEPS students with 3.0 or better gpa

```
SELECT sname FROM student
```

```
WHERE gpa >= 3.0
```

```
AND major IN ( SELECT dname FROM department  
                WHERE college = 'ceps')
```

*subquery evaluates to a set of departments*

$\forall$   $\theta$  is a comparison operator, i.e., =, <, ...



# Statistical Subquery

## Set Valued

- compute students' in-major gpas

```
SELECT sname, student.snbr,  
       SUM( cr * grade ) / sum( cr ) AS majgpa  
FROM ( student JOIN transcript  
       ON student.snbr = transcript.snbr AND major = dept  
       )  
   NATURAL JOIN course  
GROUP BY sname, student.snbr;
```

- department with highest average gpa  
– multi-valued subquery

```
SELECT major, avg( gpa ) FROM student  
GROUP BY major  
HAVING avg( gpa ) >= ALL  
  ( SELECT AVG(gpa) FROM student  
    GROUP BY major  
  );
```

# Existence Subquery Result Usage

preceded by an EXISTS or NOT EXISTS quantifier  
determines whether a nested subquery qualifies any rows  
row count is all that matters, not values  
select attribute (often \*) since actual values are not pulled out

## EXISTS

succeeds when inner query produces non-empty result

- at least one row

fails when it produces no rows — empty result

## NOT EXISTS

succeeds when inner query produces no rows

- empty table with 0 rows

fails when it produces non-empty table

# Subquery Evaluation

query evaluation proceeds inside out  
like parenthesized subexpressions

subquery evaluation is either correlated or uncorrelated  
uncorrelated

inner select does not reference outer select  
inner & outer select evaluate independently

- inner select evaluates once at start
- outer select evaluates starts after inner result available
- efficient

correlated

inner select references outer select attribute, table, or alias  
inner & outer select evaluation is interdependent,  $\therefore$

- inner select evaluates once for each outer select row
- powerful, but can be inefficient

# Uncorrelated Multi-valued Subquery

CEPS students with 3.0 or better gpa – three variations

```
SELECT sname FROM student
WHERE gpa >= 3.0
      AND major IN ( SELECT dname FROM department WHERE college = 'science' )
```

```
SELECT sname FROM student
WHERE gpa >= 3.0
      AND major = SOME ( SELECT dname FROM department WHERE college =
                          'science' )
```

```
SELECT sname FROM student
WHERE gpa >= 3.0
      AND major = ANY ( SELECT dname FROM department WHERE college = 'science'
                        )
```

student with the highest gpa

```
SELECT sname FROM student
WHERE gpa >= ALL ( SELECT gpa FROM student)
```

# Uncorrelated Multi-valued Subqueries

- departments with no majors — two variations

```
SELECT dname FROM department
WHERE dname NOT IN ( SELECT major FROM student )
```

```
SELECT dname FROM department
WHERE dname <> ALL ( SELECT major FROM student )
```

- list departments with majors — two variations

```
SELECT dname FROM department
WHERE dname IN ( SELECT major FROM student )
```

```
SELECT dname FROM department
WHERE dname = ANY ( SELECT major FROM student )
```

- cs students whose advisor is from outside 'cs'

```
SELECT sname FROM student
WHERE major = 'cs' AND advisor NOT IN
      ( SELECT name FROM faculty WHERE dept = 'cs' )
```

*join is a better way for this*

# Correlated Multi-valued Subquery

- students whose advisors are from their respective majors

```
SELECT sname FROM student WHERE advisor IN  
  ( SELECT name FROM faculty WHERE dname = major )
```

```
SELECT sname FROM student WHERE advisor =  
  ANY ( SELECT name FROM faculty WHERE dname = major )
```

- subquery references attribute major from the outer query
- the appropriate department to check varies with each student

- students whose advisor is from outside their major

```
SELECT sname FROM student WHERE advisor NOT IN  
  ( SELECT name FROM faculty WHERE dname = major )
```

```
SELECT sname FROM student WHERE advisor <> ALL  
  ( SELECT name FROM faculty WHERE dname = major )
```

- inefficient, subquery recomputes department faculty for each major

# Correlated Multi-valued Subquery - Aliases

- students taking at least one course in their respective major

```
SELECT sname FROM student S
WHERE S.major IN
      ( SELECT dept FROM enroll WHERE snbr = S.snbr )
```

– snbr refers to table enroll

S.snbr refers to table student

- student(s) in each major with the highest gpa in the major

```
SELECT sname FROM student S
WHERE gpa >= ALL
      ( SELECT gpa FROM student WHERE major = S.major )
```

inner & outer both range over same table

need alias to distinguish references

unqualified major refers to inner subquery

qualified major refers to the outer select

# Existence Subquery Patterns

SELECT subject

WHERE EXISTS ( SELECT case(s) that qualify subject )

- students taking at least 1 major course (used earlier in join)

SELECT sname FROM student WHERE EXISTS

( SELECT \* FROM enroll WHERE snbr = student.snbr AND dept =  
major );

SELECT subject

WHERE NOT EXISTS ( SELECT case(s) that disqualify subject )

- students taking no major courses

SELECT sname FROM student WHERE NOT EXISTS

( SELECT \* FROM enroll WHERE snbr = student.snbr AND dept =  
major );

- students taking only major courses (no non-major courses)

SELECT sname FROM student WHERE NOT EXISTS

( SELECT \* FROM enroll WHERE snbr = student.snbr AND dept <>  
major );

- students not enrolled in any courses

SELECT sname FROM student WHERE NOT EXISTS

( SELECT \* FROM enroll WHERE snbr = student.snbr )



# Nested Existence Subqueries

SELECT subject (that satisfies all requirements)  
WHERE NOT EXISTS

SELECT requirements the subject must meet  
WHERE NOT EXISTS

SELECT case where subject fails the requirement

principle: double negative  $\Rightarrow$  positive

if there is no failed prerequisite, all prerequisites are met

Example: course enrollments where student satisfies all  
prereqs -- note pattern: ( .. not exists .. ( not exists .. ) )

SELECT sname, dept, cnbr FROM student S JOIN enroll E ON S.snbr = E.snbr  
WHERE NOT EXISTS

*succeeds when student has*

( SELECT \* FROM prereq  
WHERE dept = E.dept  
AND cnbr = E.cnbr  
AND NOT EXISTS

*completed all prerequisites*

( SELECT \* FROM transcript  
WHERE snbr = S.snbr

*succeeds when student has*

*not completed a prerequisite*

*fails if a prereq is*

*completed*

AND dept = pdept AND cnbr = pcnbr  
AND grade >= 1.0  
) );

# Similar Existence Queries

- students who satisfy all prereqs for all enrollments

```
SELECT sname FROM student S
WHERE NOT EXISTS
  ( SELECT * FROM enroll E
      JOIN prereq P ON P.dept = E.dept AND P.cnbr = E.cnbr
    WHERE E.snbr = S.snbr AND NOT EXISTS
      ( SELECT * FROM transcript
        WHERE snbr = S.snbr AND grade >= 1.0
        AND dept = pdept AND cnbr = pcnbr
      )
    );
```

- students who satisfy their degree plan (same structure)

```
SELECT      snbr, sname FROM student S
WHERE      NOT EXISTS
  ( SELECT * FROM degrplan D
    WHERE id = S.snbr AND NOT EXISTS
      ( SELECT * FROM transcript
        WHERE snbr = D.id AND grade >= 1.0
        AND dept = D.dname AND cnbr = D.cnbr
      )
    );
```

# Existence Subqueries — ( exists ( not exists ))

- enrollments where student skipped a prereq

```
SELECT sname, dept, cnbr
FROM student JOIN enroll ON student.snbr = enroll.snbr
WHERE EXISTS
    ( SELECT * FROM prereq
      WHERE dept = enroll.dept
        AND cnbr = enroll.cnbr
        AND NOT EXISTS
            ( SELECT * FROM transcript
              WHERE snbr = student.snbr
                AND dept = pdept
                AND cnbr = pcnbr
                AND grade >= 1.0
            )
    )
);
```

*succeeds when student has  
not completed a prerequisite*

*succeeds when student has  
not completed a prerequisite  
fails if a prereq is completed*

# Existence Subqueries — ( not exists )

- enrollments & prereq that student skipped

```
SELECT sname, enroll.dept, enroll.cnbr, pdept, pcnbr
FROM student
      JOIN enroll ON student.snbr = enroll.snbr
      JOIN prereq ON enroll.dept = prereq.dept
                  AND enroll.cnbr = prereq.cnbr
WHERE  NOT EXISTS
      ( SELECT * FROM transcript
        WHERE snbr = student.snbr
          AND dept = pdept
          AND cnbr = pcnbr
          AND grade >= 1.0
      );
```

# Existence Subqueries – ( not exists ( exists ))

- enrollments where student satisfies no prerequisites

```
SELECT sname, dept, cnbr
FROM student
      JOIN enroll ON student.snbr = enroll.snbr
WHERE NOT EXISTS
( SELECT * FROM prereq
  WHERE dept = enroll.dept
    AND cnbr = enroll.cnbr
    AND EXISTS
      ( SELECT * FROM transcript
        WHERE snbr = student.snbr
          AND dept = pdept
          AND cnbr = pcnbr
          AND grade >= 1.0
      )
);
```

*succeeds when student has  
not completed a prerequisite*

*succeeds when student has  
completed a prerequisite  
fails if a prereq is uncompleted*

# Existence Subqueries — four levels

- students who satisfy all prerequisites for all enrollments

▪ ... ( NOT EXISTS ( EXISTS ( NOT EXISTS ... ) ) ) ...

```
SELECT sname
FROM student S
WHERE NOT EXISTS
  ( SELECT * FROM student
    JOIN enroll ON student.id = enroll.id
    WHERE student.id = S.id
      AND EXISTS
        ( SELECT * FROM prereq
          WHERE dept = enroll.dept
            AND cnbr = enroll.cnbr
            AND NOT EXISTS
              ( SELECT * FROM transcript
                WHERE id = student.id
                  AND dept = pdept
                  AND cnbr = pcnbr
                  AND grade >= 1.0
              )
        )
      )
  )
);
```

*fails when student has an enrollment but has not satisfied some prereq*

*succeeds when student has not completed all prereq*

*fails when student has completed the prereq*

*succeeds otherwise*

# Subquery vs. Join

- use the more intuitive form
  - i.e., which syntax better conveys the nature of the query
  - join often more intuitive for an existing relationship
    - Like a FK/PK
  - subquery often more intuitive for
    - creating an ad-hoc relationship
    - passing an aggregate value to the outer query
    - isolating parts of a long complex query from each other
- inner joins often more efficient than correlated subqueries
  - most vendors make considerable effort to optimize joins
  - especially true for defined relationships such as PK/FK
- result columns must come from table(s) in the outer query
  - join often helps put more tables in the outer query

# Subquery vs. Join: Obvious Relationship

- join often more intuitive for existing relationships (FK/PK)
  - query to list names of students enrolled in cs courses

SELECT sname	SELECT sname
FROM student S	FROM student S
JOIN enroll E	WHERE S.snbr IN
ON E.snbr = S.snbr	( SELECT snbr
WHERE E.dept = 'cs';	FROM enroll
	WHERE dept = 'cs' );

- query to list students and their teachers this semester

SELECT DISTINCT sname, name	SELECT DISTINCT sname, name
FROM student S	FROM student S, instr E
JOIN enroll E	WHERE EXISTS
ON S.snbr = E.snbr	( SELECT * FROM enroll
JOIN instr I	WHERE snbr = S.snbr
ON E.dept = I.dname	AND dept = I.dname
AND E.cnbr = I.cnbr	AND cnbr = I.cnbr
AND E.sect = I.sect;	AND sect = I.sect );



# Subquery vs. Join: Statistical Summary

- subquery value for the outer query

```
SELECT sname FROM student      join version ???  
WHERE gpa >=  
      ( SELECT AVG( gpa ) FROM student );
```

- subquery for a complex query

```
SELECT * FROM student S      join version ???  
WHERE NOT EXISTS  
      ( SELECT * FROM enroll E  
        WHERE snbr = S.snbr AND EXISTS  
              ( SELECT * FROM prereq P  
                WHERE dept = E.dept AND cnbr = E.cnbr  
                  AND NOT EXISTS  
                        ( SELECT * FROM transcript  
                          WHERE dept = P.pdept AND cnbr = P.pcnbr  
                            AND snbr = S.snbr AND grade >= 1.0  
                        )  
              )  
        )  
      AND EXISTS  
      ( SELECT * FROM enroll WHERE snbr = S.snbr );
```

# Bottom Line

- if straightforward join expression is available,  
it is probably better than a subquery
  - more intuitive
  - more efficient
- otherwise, subquery expression is probably more intuitive